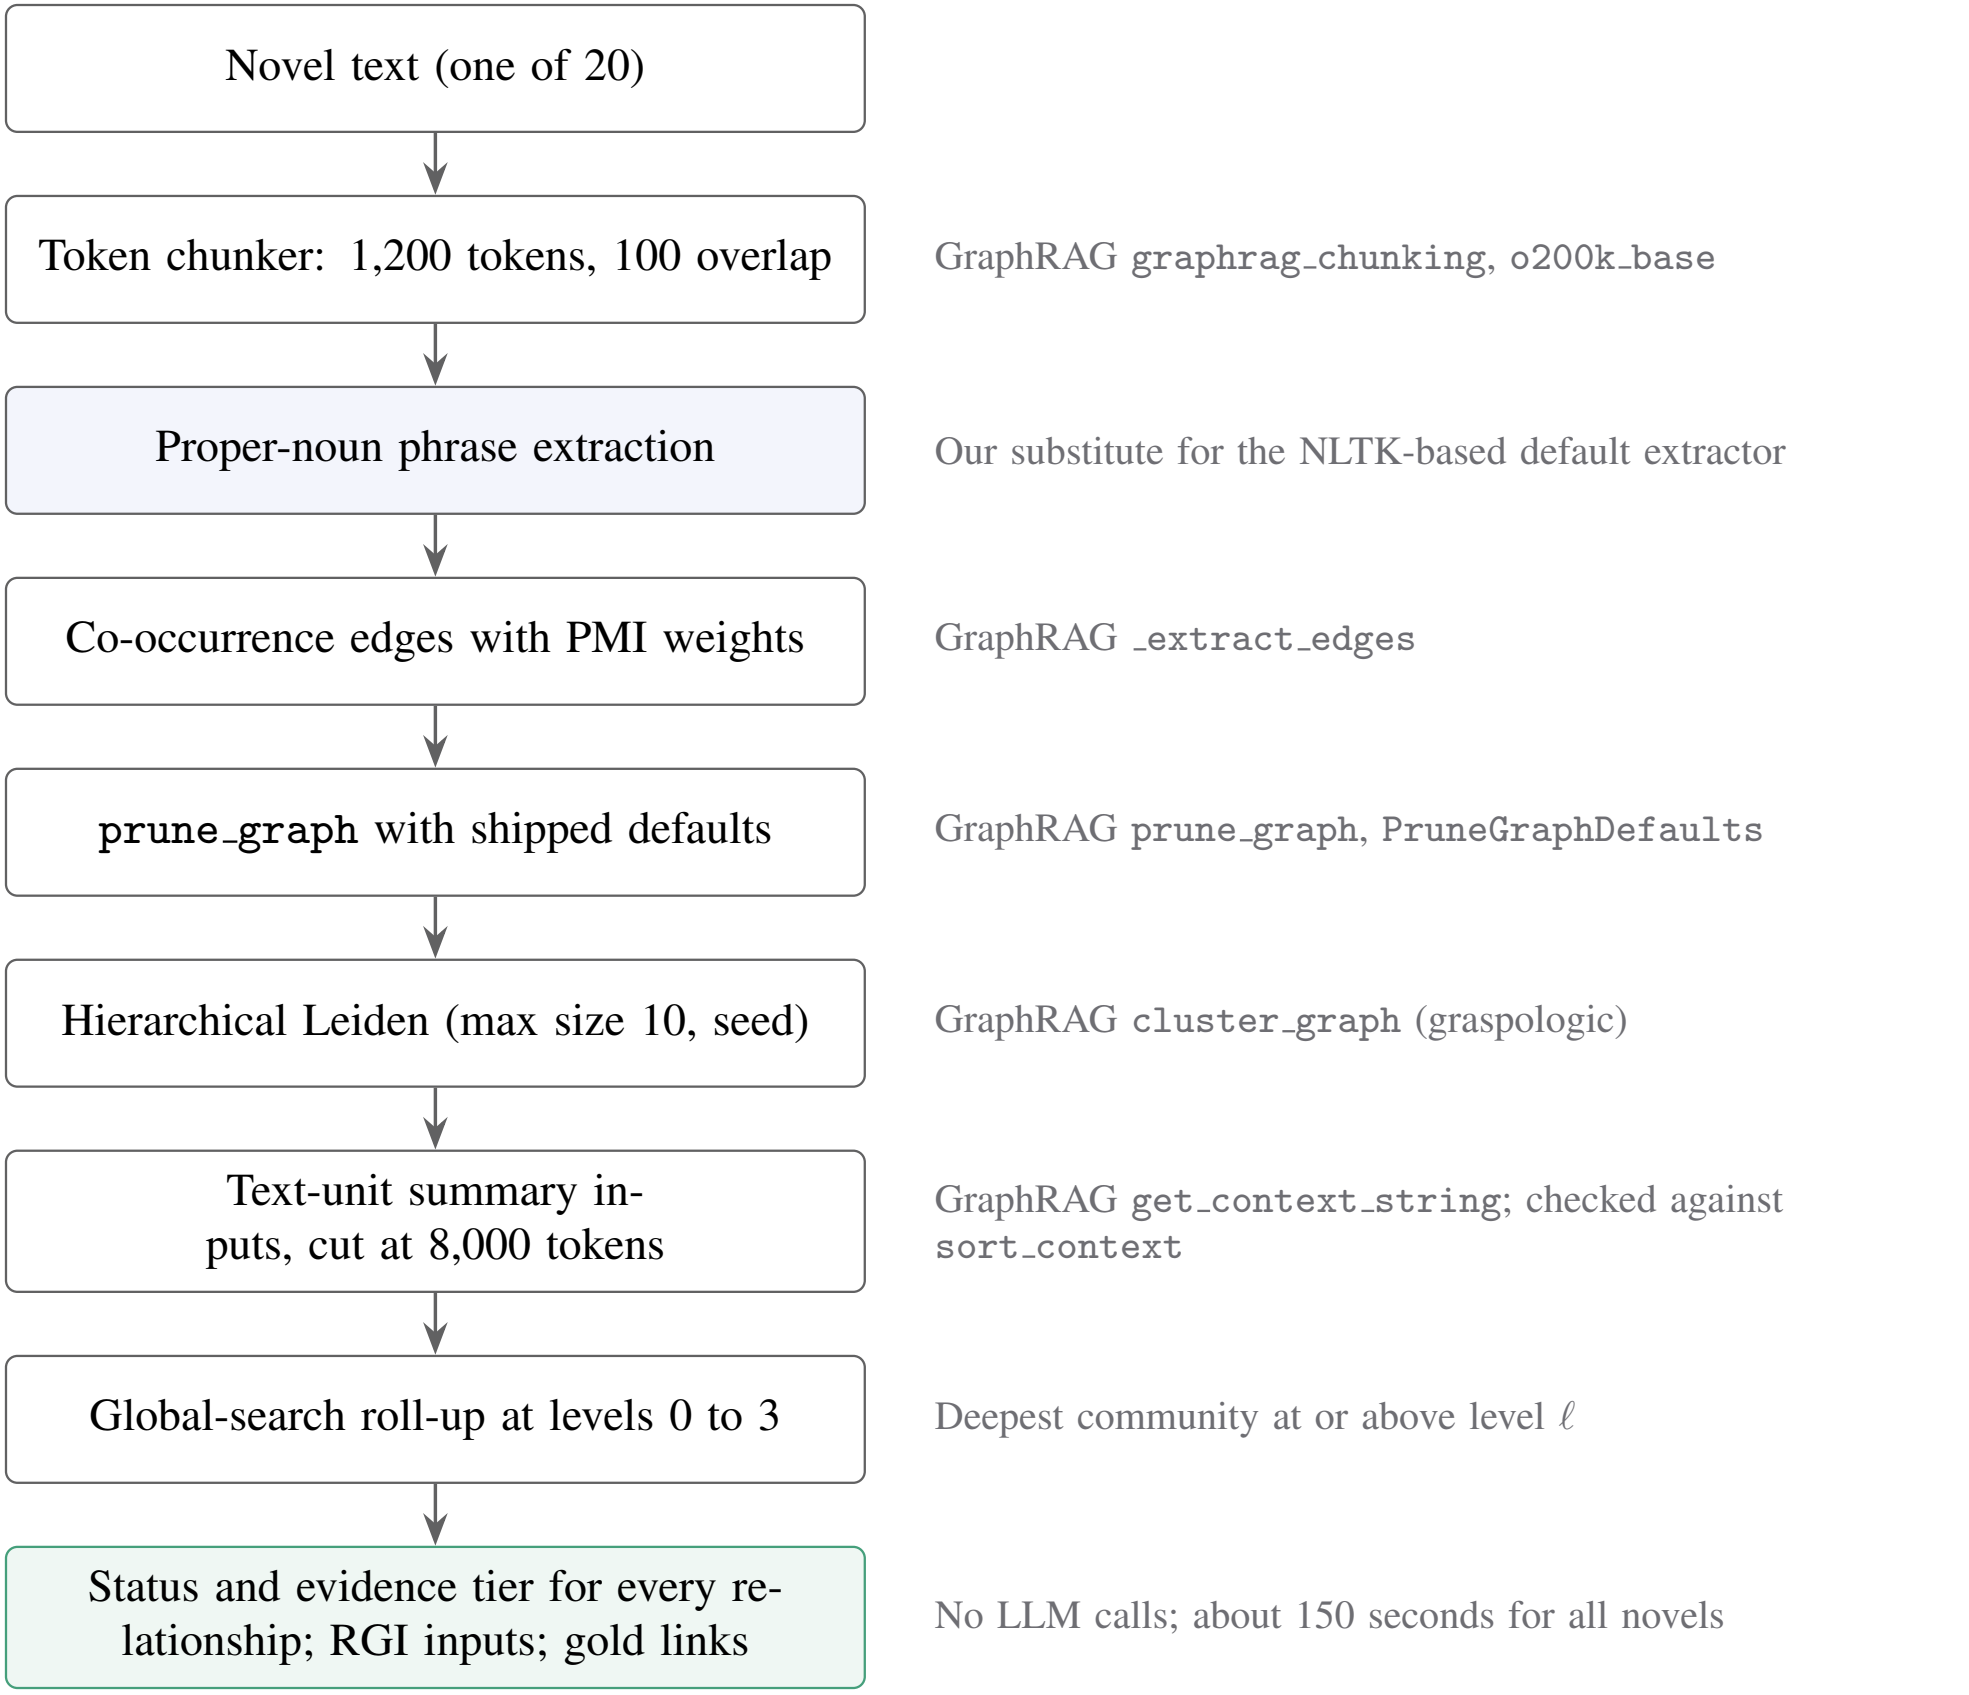
Novel text (one of 20)
Token chunker: 1,200 tokens, 100 overlap
GraphRAG graphrag_chunking, o200k_base
Proper-noun phrase extraction
Our substitute for the NLTK-based default extractor
Co-occurrence edges with PMI weights
GraphRAG _extract_edges
prune_graph with shipped defaults
GraphRAG prune_graph, PruneGraphDefaults
Hierarchical Leiden (max size 10, seed)
GraphRAG cluster_graph (graspologic)
Text-unit summary inputs, cut at 8,000 tokens
GraphRAG get_context_string; checked against sort_context
Global-search roll-up at levels 0 to 3
Deepest community at or above level ℓ
Status and evidence tier for every relationship; RGI inputs; gold links
No LLM calls; about 150 seconds for all novels